

We will be attempting to create a fraud detection model.

```
import pandas as pd

df = pd.read_csv('/creditcard.csv')
```

```
df.head()
```

```
↗
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458

5 rows x 31 columns

"The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

It contains only numerical input variables which are the result of a PCA transformation. Unfortunately, due to confidentiality issues, we cannot provide the original features and more background information about the data. Features V1, V2, ... V28 are the principal components obtained with PCA, the only features which have not been transformed with PCA are 'Time' and 'Amount'. Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset. The feature 'Amount' is the transaction Amount, this feature can be used for example-dependant cost-sensitive learning. Feature 'Class' is the response variable and it takes value 1 in case of fraud and 0 otherwise."

```
df['Class'].value_counts()
```

```
↗
```

	count
Class	
0	284315
1	492

dtype: int64

Due to the unbalanced nature of this dataset, we will have to find some work arounds because if we predicted not fraud for everything, as in the naive model, we would have a classifier with accuracy of 99.8% which would be very misleading and useless for purposes of fraud detection.

```
!pip install imbalanced-learn
```

```
↗ Requirement already satisfied: imbalanced-learn in /usr/local/lib/python3.11/dist-packages (0.13.0)
Requirement already satisfied: numpy<3,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (2.0.2)
Requirement already satisfied: scipy<2,>=1.10.1 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (1.15.2)
Requirement already satisfied: scikit-learn<2,>=1.3.2 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (1.5.2)
Requirement already satisfied: sklearn-compat<1,>=0.1 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (0.1.2)
Requirement already satisfied: joblib<2,>=1.1.1 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (1.4.2)
Requirement already satisfied: threadpoolctl<4,>=2.0.0 in /usr/local/lib/python3.11/dist-packages (from imbalanced-learn) (3.3.0)
```

```
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import SMOTE
```

```
df.dropna(inplace=True)
df.reset_index(drop=True, inplace=True)
X = df.drop('Class', axis=1)
y = df['Class']
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
# We'll only balance the training set
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.3, random_state=42, stratify=y
)
# Initialize SMOTE
sm = SMOTE(random_state=42)

# Apply SMOTE to the training set
X_resampled, y_resampled = sm.fit_resample(X_train, y_train)

print("Before SMOTE:")
print(y_train.value_counts())

print("\nAfter SMOTE:")
print(pd.Series(y_resampled).value_counts())

print("Class distribution in test set:")
print(y_test.value_counts())
```

```
⇒ Before SMOTE:
Class
0    199020
1      344
Name: count, dtype: int64

After SMOTE:
Class
0    199020
1    199020
Name: count, dtype: int64
Class distribution in test set:
Class
0    85295
1     148
Name: count, dtype: int64
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC

from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report, confusion_matrix
import numpy as np
```

```
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000)#,
    #"Random Forest": RandomForestClassifier(
    #n_estimators=50,
    #max_depth=10,
    #max_features='sqrt',
    #n_jobs=-1,
    #random_state=42
    #),
    #"SVM": SVC(kernel='rbf', probability=True)
}
for name, model in models.items():
    print(f"\n{name}")
    f1 = cross_val_score(model, X_resampled, y_resampled, cv=5, scoring='f1')
    print(f"F1 score (5-fold): {f1.mean():.4f} ± {f1.std():.4f}")

for name, model in models.items():
    model.fit(X_resampled, y_resampled)
    y_pred = model.predict(X_test)

    print(f"\n{name} - Test Set Evaluation")
    print(confusion_matrix(y_test, y_pred))
    print(classification_report(y_test, y_pred, digits=4))
```

```
⇒ Logistic Regression
F1 score (5-fold): 0.9585 ± 0.0008
```

```
Logistic Regression - Test Set Evaluation
[[83403 1892]
 [ 18 130]]
```

	precision	recall	f1-score	support
0	0.9998	0.9778	0.9887	85295
1	0.0643	0.8784	0.1198	148
accuracy			0.9776	85443
macro avg	0.5320	0.9281	0.5542	85443
weighted avg	0.9982	0.9776	0.9872	85443

```
#random forest
rf = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)
rf.fit(X_resampled, y_resampled)
```

```
↗
▼ RandomForestClassifier ⓘ ?
RandomForestClassifier(max_depth=10, random_state=42)
```

```
y_pred_rf = rf.predict(X_test)
```

```
print("matrix confusion")
print(confusion_matrix(y_test, y_pred_rf))
print("\n report")
print(classification_report(y_test, y_pred_rf, digits=4))
print("accuracy - ", accuracy_score(y_test, y_pred_rf))
```

```
↗ matrix confusion
[[85159 136]
 [ 28 120]]
```


report	precision	recall	f1-score	support
0	0.9997	0.9984	0.9990	85295
1	0.4688	0.8108	0.5941	148
accuracy			0.9981	85443
macro avg	0.7342	0.9046	0.7965	85443
weighted avg	0.9988	0.9981	0.9983	85443

accuracy - 0.9980805917395222

```
#feature importance
importances = rf.feature_importances_
features = X.columns
indices = np.argsort(importances)[::-1]
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
sns.barplot(x=importances[indices][:10], y=features[indices[:10]])
plt.title("Top 10 Feature Importances")
plt.show()
```




We analyzed feature importance using a Random Forest model to understand which variables contributed most to fraud detection model. We should keep in mind that the dataset's features are anonymized PCA components (V1 to V28). From this bar chart, it is evident that the model consistently ranked V14, V12, and V3 as the most influential features. Interpreting feature importance is important since it helps in the validation of a model's behavior. In real-world scenarios these analysis can help with building alert systems because we get to focus on a small set of variables which have the most impact.

```
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, roc_auc_score
```

```
y_probability = rf.predict_proba(X_test)[: , 1]
fpr, tpr, _ = roc_curve(y_test, y_probability)
auc = roc_auc_score(y_test, y_probability)
```

```
plt.plot(fpr, tpr, label=f'AUC = {auc}')
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("ROC Curve generated for a Random Forest model")
plt.legend()
plt.grid(True)
plt.show()
```



This is an evaluation of the model using the ROC curve, which visualizes the trade-off between true positive and false positive rates given a range of thresholds. This Random Forest model achieved an AUC of approximately 0.995, which shows a strong discriminatory force

between the fraudulent and non-fraudulent transactions.

```
import numpy as np
import plotly.graph_objs as go

# Assume y_probability, y_test, fpr, tpr, auc are already defined

# Create thresholds for the slider
thresholds = np.linspace(0, 1, 101)

# Calculate TPR and FPR for each threshold
tpr_list = []
fpr_list = []
for thresh in thresholds:
    y_pred = (y_probability >= thresh).astype(int)
    tp = np.sum((y_test == 1) & (y_pred == 1))
    fp = np.sum((y_test == 0) & (y_pred == 1))
    fn = np.sum((y_test == 1) & (y_pred == 0))
    tn = np.sum((y_test == 0) & (y_pred == 0))
    tpr_val = tp / (tp + fn) if (tp + fn) > 0 else 0
    fpr_val = fp / (fp + tn) if (fp + tn) > 0 else 0
    tpr_list.append(tpr_val)
    fpr_list.append(fpr_val)

# Create the ROC curve trace
roc_trace = go.Scatter(
    x=fpr,
    y=tpr,
    mode='lines',
    name=f'AUC = {auc:.2f}'
)

# Add a marker that moves with the slider
marker_trace = go.Scatter(
    x=[fpr_list[0]],
    y=[tpr_list[0]],
    mode='markers',
    marker=dict(color='red', size=12),
    name='Threshold Point'
)


# Create steps for the slider
steps = []
for i, thresh in enumerate(thresholds):
    step = dict(
        method="update",
        args=[{"x": [fpr, [fpr_list[i]]], "y": [tpr, [tpr_list[i]]}],
        label=f"{thresh:.2f}"
    )
    steps.append(step)

sliders = [dict(
    active=0,
    currentvalue={"prefix": "Threshold: "},
    pad={"t": 50},
    steps=steps
)]

layout = go.Layout(
    title='ROC Curve with Threshold Slider',
    xaxis=dict(title='FPR', range=[-0.1, 1]),
    yaxis=dict(title='TPR', range=[0, 1.1], scaleanchor=None), # range and scaleanchor
    showlegend=True,
    sliders=sliders
)

fig = go.Figure(data=[roc_trace, marker_trace], layout=layout)
fig.update_layout(height=700)
fig.show()

sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d', cmap='Blues', xticklabels=['Legit', 'Fraud'], yticklabels=[
plt.title("Confusion Matrix in a heatmap form")
plt.xlabel("Predicted vals")
plt.ylabel("Actual vals")
plt.show()
```



```
#decision tree build
from sklearn.tree import DecisionTreeClassifier, plot_tree

dt = DecisionTreeClassifier(max_depth=5, random_state=42)
dt.fit(X_resampled, y_resampled)
y_pred_decision_tree = dt.predict(X_test)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-3-573386433d61> in <cell line: 0>()
      3
      4 dt = DecisionTreeClassifier(max_depth=5, random_state=42)
----> 5 dt.fit(X_resampled, y_resampled)
      6 y_pred_decision_tree = dt.predict(X_test)

NameError: name 'X_resampled' is not defined
```

```
print("matrix confusion")
print(confusion_matrix(y_test, y_pred_decision_tree))
print("\n report")
print(classification_report(y_test, y_pred_decision_tree, digits=4))
print("accuracy - ", accuracy_score(y_test, y_pred_decision_tree))
```

```
matrix confusion
[[82774 2521]
 [ 24 124]]

report
```

	precision	recall	f1-score	support
0	0.9997	0.9704	0.9849	85295
1	0.0469	0.8378	0.0888	148
accuracy			0.9702	85443
macro avg	0.5233	0.9041	0.5368	85443
weighted avg	0.9981	0.9702	0.9833	85443

accuracy - 0.9702140608358789

```
#roc for decision tree
y_probability_decision = dt.predict_proba(X_test)[: , 1]
fpr_dtrees, tpr_dtrees, _ = roc_curve(y_test, y_probability_decision)
auc_dtrees = roc_auc_score(y_test, y_probability_decision)
```

```

NameError                                Traceback (most recent call last)
<ipython-input-2-cd9910676264> in <cell line: 0>()
      1 #roc for decision tree
----> 2 y_probability_decision = dt.predict_proba(X_test)[: , 1]
      3 fpr_dtree, tpr_dtree, _ = roc_curve(y_test, y_probability_decision)
      4 auc_dtree = roc_auc_score(y_test, y_probability_decision)


NameError: name 'dt' is not defined

```

```

plt.plot(fpr_dtree, tpr_dtree, label=f"Decision Tree AUC = {auc_dtree}")
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("ROC Curve of a Decision Tree")
plt.legend()
plt.grid(True)
plt.show()

```



```

import numpy as np
import plotly.graph_objs as go

# Create thresholds for the slider
thresholds_dtree = np.linspace(0, 1, 101)

# Calculate TPR and FPR for each threshold for the decision tree
fpr_list_dtree = []
tpr_list_dtree = []
for thresh in thresholds_dtree:
    y_pred_thresh = (y_probability_decision >= thresh).astype(int)
    tp = np.sum((y_test == 1) & (y_pred_thresh == 1))
    fp = np.sum((y_test == 0) & (y_pred_thresh == 1))
    fn = np.sum((y_test == 1) & (y_pred_thresh == 0))
    tn = np.sum((y_test == 0) & (y_pred_thresh == 0))
    tpr_val = tp / (tp + fn) if (tp + fn) > 0 else 0
    fpr_val = fp / (fp + tn) if (fp + tn) > 0 else 0
    tpr_list_dtree.append(tpr_val)
    fpr_list_dtree.append(fpr_val)

# Create the ROC curve trace for the decision tree
roc_trace_dtree = go.Scatter(
    x=fpr_dtree,
    y=tpr_dtree,
    mode='lines',
    name=f'Decision Tree AUC = {auc_dtree:.3f}'
)

# Add a marker that moves with the slider
marker_trace_dtree = go.Scatter(
    x=[fpr_list_dtree[0]],
    y=[tpr_list_dtree[0]],

```

```

        mode='markers',
        marker=dict(color='red', size=12),
        name='Threshold Point'
    )

# Create steps for the slider
steps_dtree = []
for i, thresh in enumerate(thresholds_dtree):
    step = dict(
        method="update",
        args=[{"x": [fpr_dtree, [fpr_list_dtree[i]]], "y": [tpr_dtree, [tpr_list_dtree[i]]}],
        label=f"{thresh:.2f}"
    )
    steps_dtree.append(step)

sliders_dtree = [dict(
    active=0,
    currentvalue={"prefix": "Threshold: "},
    pad={"t": 50},
    steps=steps_dtree
)]

layout_dtree = go.Layout(
    title=f'ROC Curve of a Decision Tree (AUC = {auc_dtree:.3f})',
    xaxis=dict(title='FPR', range=[0, 1]),
    yaxis=dict(title='TPR', range=[0, 1.1]),
    showlegend=True,
    sliders=sliders_dtree
)

fig_dtree = go.Figure(data=[roc_trace_dtree, marker_trace_dtree], layout=layout_dtree)
fig_dtree.update_layout(height=700)
fig_dtree.show()

```



```

-----
NameError                                Traceback (most recent call last)
<ipython-input-1-a970cd93321c> in <cell line: 0>()
      9 tpr_list_dtree = []
     10 for thresh in thresholds_dtree:
----> 11     y_pred_thresh = (y_probability_decision >= thresh).astype(int)
     12     tp = np.sum((y_test == 1) & (y_pred_thresh == 1))
     13     fp = np.sum((y_test == 0) & (y_pred_thresh == 1))

NameError: name 'y_probability_decision' is not defined

```

```


plt.figure(figsize=(20,10))
plot_tree(dt, feature_names=X.columns, class_names=['Legit', 'Fraud'], filled=True, max_depth=3)
plt.title("Decision Tree Visualization")
plt.show()

```





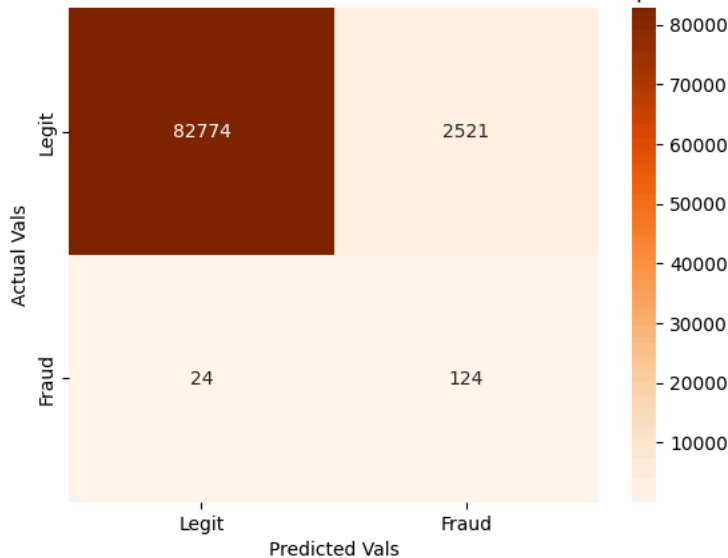
## Decision Tree Visualization



```
sns.heatmap(confusion_matrix(y_test, y_pred_decision_tree), annot=True, fmt='d', cmap='Oranges', xticklabels=['Legit', 'Fraud'],
plt.title("Confusion Matrix of a Decision Tree in form of a HeatMap")
plt.xlabel("Predicted Vals")
plt.ylabel("Actual Vals")
plt.show()
```



## Confusion Matrix of a Decision Tree in form of a HeatMap



```
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

mlp = MLPClassifier(hidden_layer_sizes=(150,50, 30),
                    activation='tanh',
                    solver='adam',
                    learning_rate_init = 0.001,
                    max_iter=150,
                    random_state=42,
                    verbose=True)

mlp.fit(X_resampled, y_resampled)

y_pred = mlp.predict(X_test)
```

```
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("\nAccuracy Score:")
print(accuracy_score(y_test, y_pred))
```

```
↩ Iteration 1, loss = 0.03367488
  Iteration 2, loss = 0.00572472
  Iteration 3, loss = 0.00348618
  Iteration 4, loss = 0.00250909
  Iteration 5, loss = 0.00209267
  Iteration 6, loss = 0.00177511
  Iteration 7, loss = 0.00152416
  Iteration 8, loss = 0.00145739
  Iteration 9, loss = 0.00102083
  Iteration 10, loss = 0.00128659
  Iteration 11, loss = 0.00102475
  Iteration 12, loss = 0.00118519
  Iteration 13, loss = 0.00096564
  Iteration 14, loss = 0.00091743
  Iteration 15, loss = 0.00097527
  Iteration 16, loss = 0.00080420
  Iteration 17, loss = 0.00102345
  Iteration 18, loss = 0.00085224
  Iteration 19, loss = 0.00095319
  Iteration 20, loss = 0.00077160
  Iteration 21, loss = 0.00083653
  Iteration 22, loss = 0.00079524
  Iteration 23, loss = 0.00064409
  Iteration 24, loss = 0.00068423
  Iteration 25, loss = 0.00075880
  Iteration 26, loss = 0.00079214
  Iteration 27, loss = 0.00071163
  Iteration 28, loss = 0.00077091
  Iteration 29, loss = 0.00065425
  Iteration 30, loss = 0.00063837
  Iteration 31, loss = 0.00062175
  Iteration 32, loss = 0.00067103
  Iteration 33, loss = 0.00068373
  Iteration 34, loss = 0.00065193
Training loss did not improve more than tol=0.000100 for 10 consecutive epochs. Stopping.
Confusion Matrix:
[[85256   39]
 [   36  112]]

Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	85295
1	0.74	0.76	0.75	148
accuracy			1.00	85443
macro avg	0.87	0.88	0.87	85443
weighted avg	1.00	1.00	1.00	85443

```

Accuracy Score:
0.9991222218320986

```

```
#Grid Search to find optimal hidden layer sizes
from sklearn.model_selection import GridSearchCV
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import make_scorer, f1_score

param_grid = {
    'hidden_layer_sizes': [
        (30,), (50,), (150,),
        (50, 30), (150, 50), (150, 50, 30)
    ]
}

mlp = MLPClassifier(
    activation='tanh',
    solver='adam',
    learning_rate_init = 0.001,
    max_iter=150,
    random_state=42,
)
```

```
scorer = make_scorer(f1_score, pos_label=1) # Use F1-score for the minority class (1) as the evaluation metri
```

```
grid_search = GridSearchCV(  
    estimator=mlp,  
    param_grid=param_grid,  
    scoring=scorer,  
    cv=5,  
    verbose=2,  
    n_jobs=-1  
)
```

```
grid_search.fit(X_resampled, y_resampled)
```

```
print("Best hidden_layer_sizes:", grid_search.best_params_)  
print("Best F1 score (minority class):". grid_search.best_score )
```

```
➡ Fitting 5 folds for each of 6 candidates, totalling 30 fits  
Best hidden_layer_sizes: {'hidden_layer_sizes': (150, 50, 30)}  
Best F1 score (minority class): 0.9997413005950356
```

```
from sklearn.metrics import roc_curve, roc_auc_score  
import matplotlib.pyplot as plt
```

```
y_proba = mlp.predict_proba(X_test)[:, 1]
```

```
fpr, tpr, thresholds = roc_curve(y_test, y_proba)  
auc_score = roc_auc_score(y_test, y_proba)
```

```
plt.figure(figsize=(8, 6))  
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {auc_score:.4f})')  
plt.plot([0, 1], [0, 1], 'k--', label='Naive Classifier')  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
plt.title('NN Model ROC Curve')  
plt.legend(loc='lower right')
```